

ECEN 4532: Digital Signal Processing Lab

Lecture Notes: Lab 3

Instructor: Prof. Farhad Pourkamali-Anaraki

University of Colorado at Boulder

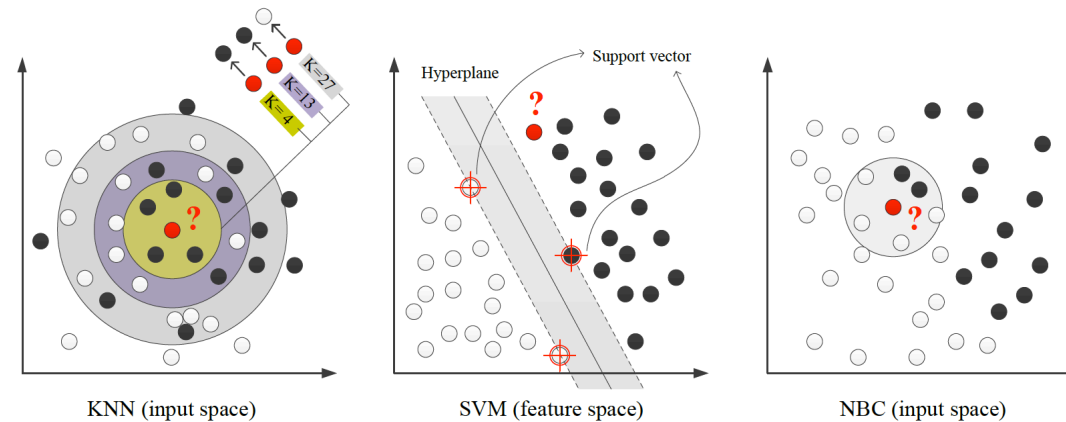
Spring 2017



University of Colorado
Boulder

Introduction

- In this lab, we will use **classification** methods to automatically detect the genre of a track.



- Database of 150 training samples (6x25)
- Please download these tracks at: (valid for 7 days)

<https://www.dropbox.com/s/rcr3ps100eu1hgg/Lab3-audio-tracks.zip?dl=0>


Introduction

- Therefore, we need to develop **distances** to compare audio tracks.
- The distances are based on the **features** that we developed in the first and second labs:
 - MFCC coefficients
 - Chroma (Normalized Pitch Class Profile)

Genre Classification

- We will consider the following genres:

1. classical
2. electronic
3. jazz/blues
4. metal/punk
5. pop/rock
6. world



Statistical Evaluation of Algorithms

- **Confusion Matrix:** The **accuracy** of the classification will be quantified using confusion matrices.
- For a given classification experiment, you will construct a **6×6 matrix** where the **rows** are the true genres, and the **columns** are the genres classified by your algorithm.

Confusion Matrices

- The entry $R(i, j)$ of the confusion matrix R is the number of songs of genre i that were classified as genre j .

		predicted genre by your algorithm					
true genre	classical	21	3	1	0	0	0
	electronics	3	22	0	0	0	0
	jazz	0	1	13	2	8	1
	punk	0	0	0	23	1	1
	rock	0	2	3	6	9	5
	world	3	7	1	0	4	10

Confusion Matrices

- Diagonal entries of the confusion matrix:

	classical	electronic	jazz	punk	rock	world
classical	21	3	1	0	0	0
electronics	3	22	0	0	0	0
jazz	0	1	13	2	8	1
punk	0	0	0	23	1	1
rock	0	2	3	6	9	5
world	3	7	1	0	4	10

- The **correct classification rate** per class was:

[0.84 0.88 0.52 0.92 0.36 0.40]

Cross-Validation

- Cross-validation is a technique for assessing how the results of a statistical analysis will **generalize** to an independent data set.
- In this lab, you have access to 25 songs per genre, so $25 \times 6 = 150$ total songs.

$$\mathcal{S} = \bigcup_{i=1}^6 \mathcal{G}_i$$

Cross-Validation

- Randomly divide each genre into 5 subsets of size 5:

$$\mathcal{G}_i = \bigcup_{k=1}^5 \mathcal{G}_i^k$$

- where: $|\mathcal{G}_i^k| = |\mathcal{G}_i|/5 = 5$

Cross-Validation

for $k = 1$ **to** 5 // round robin over the testing songs

- form the set of test songs

$$\mathcal{U} = \bigcup_{i=1}^6 \{\mathcal{G}_i^k, \}$$

and the set of training songs

$$\mathcal{L} = \bigcup_{i=1}^6 \{\mathcal{G}_i^l, l = 1, \dots, 5, l \neq k\}$$

- test the performance of your algorithm on the 30 songs in \mathcal{U} using the 120 training songs \mathcal{L}
- record the confusion matrix

Cross-Validation

- Finally, repeat this procedure 10 different times over various randomizations.
- You will compute the **mean** and **standard** deviation for all the entries in the confusion matrices
 - for each entry (i, j) of the confusion matrix:
 - compute the average $\bar{R}(i, j)$
 - compute the standard deviation $\sigma_R(i, j)$

Random Permutation in MATLAB

Description

`p = randperm(n)` returns a row vector containing a random permutation of the integers from 1 to `n` inclusive.

```
>> randperm(10)
```

```
ans =
```

```
     6     3     7     8     5     1     2     4     9    10
```

`rng(seed)` seeds the random number generator using the nonnegative integer `seed` so that `rand`, `randi`, and `randn` produce a predictable sequence of numbers.

Cross-Validation in MATLAB

crossvalind

Generate cross-validation indices

Syntax

```
Indices = crossvalind('Kfold', N, K)
[Train, Test] = crossvalind('HoldOut', N, P)
[Train, Test] = crossvalind('LeaveMOut', N, M)
[Train, Test] = crossvalind('Resubstitution', N, [P,Q])
[...] = crossvalind(Method, Group, ...)
[...] = crossvalind(Method, Group, ..., 'Classes', C)
[...] = crossvalind(Method, Group, ..., 'Min', MinValue)
```

Cross-Validation in MATLAB

```
>> N = 25; [Train, Test] = crossvalind('LeaveMOut', N, 5);
```

```
>> size(Train)
```

```
ans = 25 1
```

```
>> size(Test)
```

```
ans = 25 1
```

```
>> size(find(Train==1),1)
```

```
ans = 20
```

```
>> size(find(Test==1),1)
```

```
ans = 5
```

Property of a distance between tracks

- The distance should provide a natural **partitioning** of the dataset.
- Let X_i be the set of features that you compute for track i .

1. for a given genre g , we define its centroid as

$$\overline{\mathbf{X}}_g = \frac{1}{25} \sum_{k \in \text{genre } g} \mathbf{X}_k$$

2. We also define the radius ρ_g of a genre g as the standard deviation of the genre

$$\rho_g = \frac{1}{25} \sum_{k \in \text{genre } g} d(\mathbf{X}_k, \overline{\mathbf{X}}_g)$$

Property of a distance between tracks

3. Finally, we require that the genres do not overlap, and therefore if we consider two genres, g and g' ,

$$d(\overline{\mathbf{X}}_g, \overline{\mathbf{X}}_{g'}) > \rho_g + \rho_{g'} \quad (4)$$

Distances Between Tracks

- The key ingredient of classification methods is a **distance** that quantifies the **similarity** between tracks in terms of musical features.
- We consider the following features:
 - MFCC coefficients
 - Chroma (Normalized Pitch Class Profile)

	1	2		...					N_f
1									
\vdots									
K									

Resizing MFCC Coefficient

- In order to obtain more reliable classification results and speed up the computation, we merge some of the Mel banks:

```
t = zeros(1,36); (2.21)
t(1) =1;t(7:8)=5;t(15:18)= 9;
t(2)    = 2; t( 9:10) = 6; t(19:23) = 10;
t(3:4) = 3; t(11:12) = 7; t(24:29) = 11;
t(5:6) = 4; t(13:14) = 8; t(30:36) = 12;
```

```
mel2 = zeros(12,size(mfcc,2));
```

```
for i=1:12,
    mel2(i,:) = sum(mfcc(t==i,:),1);
end
```

Resizing MFCC Coefficient

```
>> t = [1,1,1,1,2];  
>> t == 1
```

```
ans =
```

```
     1     1     1     1     0
```

```
>> t(ans)
```

```
ans =
```

```
     1     1     1     1
```

Resizing MFCC Coefficient

```
--  
>> A = [1 3 2; 4 2 5; 6 1 4]
```

```
A =
```

```
     1     3     2  
     4     2     5  
     6     1     4
```

```
>> sum(A,2)
```

```
ans =
```

```
     6  
    11  
    11
```

```
>> sum(A,1)
```

```
ans =
```

```
    11     6    11
```

A Probabilistic Model

- For each track, the distribution of vectors $X(:, n)$, $n = 1, \dots, N_f$ is modeled as a multivariate **Gaussian** distribution in \mathbb{R}^K , $K = 12$.
- Therefore, this approach collapses all the frames together and summarizes each track by a **mean** and a **covariance matrix**.

A Probabilistic Model

```
>> mu = mean(mfcc,2);  
>> Cov = cov(mfcc');
```

A Probabilistic Model

```
>> help cov
```

```
cov Covariance matrix.
```

```
cov(X), if X is a vector, returns the variance. For matrices,  
where each row is an observation, and each column a variable,  
cov(X) is the covariance matrix. DIAG(cov(X)) is a vector of  
variances for each column, and SQRT(DIAG(cov(X))) is a vector  
of standard deviations. cov(X,Y), where X and Y are matrices with  
the same number of elements, is equivalent to cov([X(:) Y(:)]).
```

Distance in Probabilistic Model

- First, we compute the Kullback-Leibler divergence:

$$KL(G^s, G^{s'}) = \frac{1}{2} \left(\text{tr}(\Sigma_{s'}^{-1} \Sigma_s) + (\mu_{s'} - \mu_s)^T \Sigma_{s'}^{-1} (\mu_{s'} - \mu_s) - K + \log \left(\frac{\det \Sigma_{s'}}{\det \Sigma_s} \right) \right)$$

Symmetric KL Divergence

$$KL(G^s, G^{s'}) = \frac{1}{2} \text{tr}(\Sigma_{s'}^{-1} \Sigma_s + \Sigma_s^{-1} \Sigma_{s'}) - K + \frac{1}{2} (\mu_s - \mu_{s'})^T (\Sigma_{s'}^{-1} + \Sigma_s^{-1}) (\mu_s - \mu_{s'})$$

Distance in Probabilistic Model

- Second, the KL distance is rescaled using an exponential kernel, and we define the distance between two tracks as follows:

$$d(s, s') = \exp \left(-\gamma KL(G^s, G^{s'}) \right)$$

- In this kernel, the parameter γ is chosen in $[0,1]$ to optimize the classification.

Distance Matrix

- We Compute the 6x6 average distance matrix between the genres, defined by:

$$\overline{D}(i, j) = \frac{1}{25^2} \sum_{s \in \text{genre } i, s' \in \text{genre } j} d(s, s'), \quad i, j = 1, \dots, 6.$$

- Now, you find a value of γ that maximizes the separation between the different genres.

Appendix 1: Singular Value Decomposition (SVD)

- There exists a factorization of matrix M:

$$M = U\Sigma V^T$$


diagonal matrix

Appendix 1: Singular Value Decomposition (SVD)

- What is the application of SVD?

M =

1	2	3
0	0	0
0	0	0

[U,Sigma,V] = svd(M);

Sigma =

3.7417	0	0
0	0	0
0	0	0

U =

1.0000	0	0
0	0	1.0000
0	1.0000	0

Appendix 2: Pseudo Inverse

- The SVD can be used for computing the pseudo inverse of a matrix:

$$M^+ = V\Sigma^+U^T$$



replacing every non-zero diagonal entry by its inverse